

Design and Implementation of a Web System for Searching and Managing Courses - Course Finder

First Report

Brayan Alejandro Riveros Rodríguez

Cristian Santiago Ríos Vásquez

Carlos Andres Pescador Castro

Systems Engineering

Universidad Distrital Francisco José de Caldas

December 11, 2025

Abstract

Students often must search through numerous of courses, multiple instructors, and changing schedules across semesters, but information is often not centralized or in a palatable format, or in partially synchronized systems, making discovery and planning difficult. This project proposes a web-based system that consolidates course data and formalizes instructor-course relationships, separating authentication/authorization services from central domain-based course services, with non-functional objectives to integrate and centralize data, efficient search, and course-instructor associations. Workshops 1 and 2 established the central problem using business model canvas, user stories, CRC cards, and overall system design. Guided by Scrum as a development management methodology, preparing for sprint planning and traceability, this first report summarizes preliminary results and describes the project road map.

Contents

1	Introduction	5
1.1	Context and Motivation	5
1.2	Objectives	5
1.3	Scope	5

2	Related Work and Background	6
3	Methodology	6
3.1	Process and Team Practices	6
3.2	Scrum Implementation and Project Management	7
3.3	System Analysis and Design	7
4	System Design and Architecture	8
4.1	User Stories and Requirements	8
4.1.1	Administrative User Stories	9
4.1.2	End-User Functionality	12
4.1.3	Story Management and Traceability	13
4.2	Web User Interface	14
4.3	Architecture Overview	21
4.4	Data Model and Database Schema	23
4.4.1	Domain Entities and Relationships	23
4.4.2	Business Logic Database	24
4.4.3	Authentication Database	25
4.4.4	Schema Design Rationale	26
4.5	API Endpoints and Contracts	27
4.5.1	Authentication Service API	27
4.5.2	Business Logic Service API	29
4.5.3	API Design Principles and Error Handling	34
4.6	Deployment View	36
5	Implementation and Quality Assurance	37
5.1	Testing Strategy	37
5.1.1	Unit Testing	37
5.1.2	Acceptance Testing with Cucumber	39
5.1.3	Stress Testing with JMeter	41
5.1.4	Testing Infrastructure and Automation	42
5.2	CI/CD Pipeline	43
5.3	Code Quality and Standards	43
6	Results and Discussion	43
6.1	Testing Results	43
6.2	System Performance Evaluation	43
6.3	Achievement of Objectives	43
6.4	Limitations and Lessons Learned	43

List of Figures

1	Web User Interface (v1.0)	15
2	Web User Interface (v1.0)	16
3	Web User Interface (v1.0)	17
4	Web User Interface (v1.0)	18
5	Web User Interface (v1.0)	19
6	Web User Interface (v1.0)	20
7	Architecture Diagram (v1.0).	22
8	Class diagram showing domain entities, repositories, services, and controllers across all architectural layers.	26
9	Swagger UI documentation for the authentication service showing the login endpoint specification with request/response schemas.	28
10	Login endpoint request example showing the JSON payload structure with username and password fields.	28
11	Successful login response showing the generated JWT token that must be used for authenticated requests.	29
12	Swagger UI overview of the business logic service showing all CRUD endpoints organized by resource (professors, courses, assignments).	30
13	Professor creation endpoint documentation showing request schema with required fields (name, email, specialty) and expected response format. . .	31
14	Course CRUD endpoints documentation showing all available operations (POST, GET, PUT, DELETE) with their respective paths and authentication requirements.	32
15	Assignment creation endpoint showing request schema with professor_id, course_id, and status fields, along with the nested response structure including professor and course details.	33
16	Deployment diagram (v1.0).	36

List of Tables

1	Core User Stories for System Functionality	9
2	Schema - Academic Domain Tables	24
3	Schema - Authentication Table	25
4	API Endpoints Summary	35
5	Python Backend Unit Tests Summary	38

6	Cucumber Acceptance Test Scenarios Summary	41
---	--	----

1 Introduction

1.1 Context and Motivation

Students often have no way to organize their semester schedules. At many institutions, including our own, course information is often poorly organized. This fragmentation makes it difficult for students to discover relevant offerings of interest, as well as the professors who teach them and their schedules. The result is avoidable problems during planning and registration periods, duplicate registrations, and delays in decision-making.

To address these issues, we propose a web system for searching and managing academic courses that consolidates essential data (course codes, titles, descriptions) and formalizes the relationships between professors and courses. The system targets three primary stakeholder groups: (i) *students*, who need fast, accurate search and discovery; (ii) *professors*, who is the associated with the courses; and (iii) *administrators users*, which is responsible for loading the data and managing the application. Our current work established the problem framing and initial design assets for the development phase.

From an engineering perspective, the project is intentionally designed to practice delivering end-to-end software within a semiannual schedule by implementing an agile approach.

1.2 Objectives

- Design and specify a reliable web-based system that allows key stakeholders to search, manage and maintain accurate information about courses and instructors.
- Define a clear and feasible road map for the implementation of the project that aligns with the identified problem context, stakeholder needs, and planning.
- Synthesize the context, stakeholders, and scope of the problem into a consistent baseline that unifies terminology and guides subsequent design and development activities.

1.3 Scope

This first report defines the requirements and design boundaries of the project at its current stage. It focuses on consolidating the conceptual and architectural foundations, including problem definition, stakeholder identification, user stories, CRC cards, high-level architecture, deployment views, and early user-interface sketches. The report also introduces preliminary elements for the data model and API contracts that will guide the implementation in the next phase. This scope explicitly includes documenting the

problem context, justifying the proposed solution, and establishing an initial road map under agile methodology using Taiga as the project management tool.

2 Related Work and Background

Students need digital systems to search and manage academic information such as course offerings, instructor assignments, and scheduling. However, on many occasions they find that in fragmented data or centralized in legacy applications that lack interoperability and transparency for stakeholders. In this context, academic management systems and course catalog platforms such as Moodle, Blackboard, and Google Classroom provide partial solutions, focusing primarily on learning management rather than the structural organization of academic offerings.

From a technical point of view, modern web-based architectures rely on Restful APIs, micro-service decomposition, and standard authentication mechanisms such as JWT (JSON Web Token) to ensure secure and maintainable information exchange. These principles underpin the design of this project: a lightweight, modular system divided into an authentication service and a business This dual-back-end structure supports scalability, separation of concerns, and parallel development by different team members. Furthermore, adopting a continuous integration and deployment pipeline through tools like Docker and GitHub Actions allows the project to align with current DevOps practices that promote automation, repeatability, and early defect detection.

In terms of methodology, the use of Scrum supported support on the Taiga tool reflects an agile approach, iterative development of software engineering. Agile methods are widely used in both industry and academia to encourage adaptability, transparency, and stakeholder feedback. Previous studies and projects in software engineering have shown that integrating agile practices allows for better quality deliveries.

In summary, this work draws on principles from web-based academic management systems, requirement analysis and software design to propose a simple, fast, and verifiable solution tailored to the needs of the stakeholders.

3 Methodology

3.1 Process and Team Practices

The project will follow an agile software development approach based on the Scrum framework, adapted to the context of problem. The team uses **Taiga** as the main project

management tool to maintain the product backlog, define sprint goals, and track progress through user stories and tasks.

Roles are distributed collaboratively among the three members of the team. Carlos Andres Pescador Castro serve as the Product Owner and is the responsible for prioritizing the backlog and maintaining alignment with the defined requirements. Cristian Santiago Ríos Vásquez act as the Scrum Master, facilitating coordination among team members, ensuring adherence to Scrum events, and managing sprint reviews and retrospectives. Brayan Alejandro Riveros Rodríguez as the Development Lead, focusing on technical implementation, version control integration, ci/cd pipeline and testing coordination. All members contribute equally to documentation, gathering of requirements and continuous improvement activities.

Scrum ceremonies are conducted at a scaled frequency appropriate brief stand-up meetings are held twice per week (asynchronous when needed), sprint planning and retrospectives occur at the beginning and end of each workshop, and sprint reviews coincide with workshop submissions.

Version control and collaboration are managed through a shared Git repository, following a feature-branch workflow.

3.2 Scrum Implementation and Project Management

a

3.3 System Analysis and Design

Following the completion of the Workshops, the system analysis and design phase guide the transition from requirements specification to technical implementation. The main goal will be to ensure that the proposed architecture, domain entities, and component responsibilities are consistent with the user stories, acceptance criteria, and stakeholder needs identified during the analysis stage.

The analysis focus on consolidating the functional requirements derived from the user stories created in Workshop 1. Each user story will be mapped to a corresponding set of system components, defining clear inputs, processes, and outputs. CRC (Class–Responsibility–Collaborator) cards previously designed will serve as the foundation for assigning responsibilities among classes, while maintaining object-oriented design principles. This mapping will guarantee that each functional requirement has a direct

representation within the software model.

The system will adopt a modular architecture divided into two main backend services: an authentication and authorization service, and a core domain service for managing courses, professors, and their relationships. Both services will communicate through RESTful APIs, enabling independent deployment and scalability. Each backend will interact with its respective database—MySQL for authentication and PostgreSQL or MongoDB for the core domain—ensuring persistence and referential integrity across entities. The frontend web interface, developed with modern JavaScript framework, will consume these APIs to provide a cohesive user experience.

To formalize the design, UML diagrams are used to illustrate the system structure and interactions. A class diagram will describe the main domain entities (Professor, Course, and Link), their attributes, and associations. A deployment diagram will detail how components will be distributed across environments, showing the web client, API services, and database instances.

Non-functional requirements will also guide the design. The system will ensure security through JWT-based authentication and role-based access control, maintainability through modular coding practices, and performance through optimized database queries and API endpoints. Furthermore, logging and basic monitoring mechanisms will be defined to support future testing and evaluation.

This analysis and design process will establish a clear technical blueprint for the implementation stage to be developed in Workshop 3. By grounding each architectural decision in the previously validated requirements, the team will ensure traceability from stakeholder needs to system components and a solid foundation for the subsequent development and testing phases.

4 System Design and Architecture

4.1 User Stories and Requirements

The functional requirements of the system were captured through user stories following agile methodology principles. Each story represents a specific user need and includes well-defined acceptance criteria that guide both development and testing activities. The stories were structured using the standard template: "As a [user role], I want to [perform action], so that [achieve benefit]." This approach ensures that every feature delivers clear value to stakeholders and remains traceable throughout the development lifecycle.

Five core user stories form the foundation of the system functionality, divided into administrative capabilities, end-user search features, and role-based navigation. These stories were prioritized based on their dependency relationships and their impact on delivering value to all stakeholder groups. Table 1 provides a summary of all user stories with their dependencies and roles.

Table 1: Core User Stories for System Functionality

Story ID	Story Name	User Role	Dependencies
US_CF_01	Registration of New Professors	System Administrator	None
US_CF_02	Registration of New Courses	System Administrator	None
US_CF_03	Link Professor to Courses	System Administrator	US_CF_01, US_CF_02
US_CF_04	Search Courses by Professor and Vice Versa	System User	US_CF_03
US_CF_05	Home Page and Access Flow	Administrator /General User	US_CF_04, Auth Module

4.1.1 Administrative User Stories

Three stories address the System Administrator role, enabling the creation and management of the academic catalog that supports the search functionality.

US_CourseFinder_01: Registration of New Professors This story addresses the need to register professors in the system by capturing their personal and academic details through a dedicated registration form.

User Need: The administrator requires the ability to add new professors to the platform by entering their personal and academic information, ensuring the professor database remains current and complete.

Justification: Maintaining an up-to-date professor database ensures that users have access to accurate and complete information when consulting the system, supporting informed decision-making during course selection.

Key Validations: The system must validate all required fields before persisting data, including name, identification number, email, and specialty. Email uniqueness must be enforced to prevent duplicate registrations. The identification number must also be unique across all professor records.

Acceptance Criteria: The administrator can access the registration form from the main admin panel. All required fields are validated before saving data. The system displays a confirmation message upon successful registration. The new professor appears in the general professor list without requiring a page reload. Duplicate registrations based on email or identification number are prevented with appropriate error messages.

INVEST Characteristics: This story is independent and does not depend on other stories for development. It is negotiable in terms of specific fields and validation rules, which can be adjusted according to academic requirements. The story provides immediate value by establishing the foundational data required for the platform. Development effort is estimable based on form complexity and validation logic. The scope is small enough to complete within a single sprint. The story is testable through field validation tests, database persistence verification, and UI confirmation message checks.

US_CourseFinder_02: Registration of New Courses This story enables administrators to expand the academic catalog by adding new courses with comprehensive metadata.

User Need: The administrator wants to add new courses to the system, including details such as course name, code, description, credits, and category.

Justification: Expanding and maintaining the academic offerings ensures users have access to current course information, supporting effective academic planning and course discovery.

Key Validations: The system must validate course code uniqueness to prevent duplicate entries. All required fields must be completed before submission. If a professor is assigned during course creation, the relationship must be validated to ensure the professor exists in the database.

Acceptance Criteria: The administrator can access the course registration form from the main panel. All required fields are validated before saving. A confirmation message appears upon successful registration. The new course appears in the course list without

page reload. The system prevents creation of duplicate courses based on course code. If a professor is assigned, the relationship is correctly recorded in the database.

INVEST Characteristics: This story is independent and can be developed without dependencies on other stories. Fields and professor relationships are negotiable based on academic system needs. The story expands the information base, increasing platform utility for end users. Effort is estimable considering form implementation, validations, and database operations. The scope fits within a single sprint. Testing validates field requirements, database persistence, and relationship integrity.

US__CourseFinder__03: Link Professor to One or Multiple Courses This story implements the core relationship management functionality that establishes the many-to-many associations between professors and courses.

User Need: The administrator wants to link existing professors to one or more existing courses, establishing the teaching assignments that users will query.

Justification: Correct professor-course associations enable users to view accurate relationships when searching, supporting informed course selection based on instructor preferences.

Key Validations: Both the professor and course must exist in the database before linking. The system should prevent duplicate assignments of the same professor-course-group combination. If a course is already assigned to another professor in the same group, the system should warn the administrator or prevent the duplication based on business rules.

Acceptance Criteria: The administrator can access a dedicated interface for linking professors to courses. The system displays only existing professors and courses. The administrator can select one professor and one or multiple courses to link. Upon confirmation, the links are stored and visible in both entity profiles. Error messages appear if attempting to link non-existent entities. The system prevents or warns about duplicate course assignments. The administrator can unlink or modify existing relationships. A confirmation message appears after successful linking.

INVEST Characteristics: This story depends on US_CF_01 and US_CF_02, as both entities must exist before linking. The user interface and workflow are negotiable based on administrator feedback. The story provides core administrative control over relationships, ensuring data accuracy. Development time is estimable based on form

complexity, relationship logic, and validations. The scope fits within a single sprint for a small team. Testing validates UI interactions and backend relationship integrity.

4.1.2 End-User Functionality

Two stories address the primary use cases for general users: discovering course and professor information through search and accessing the system through role-appropriate entry points.

US_CourseFinder_04: Search Courses by Professor and Professors by Course

This story provides the core bidirectional search functionality that motivated the project development.

User Need: System users want to search for courses taught by a specific professor and search for courses to view their assigned professors.

Justification: Easy access to accurate professor-course relationships improves navigation and course discovery, helping users make informed academic decisions efficiently.

Key Validations: Search queries must handle partial text matches and ignore case sensitivity. The system must validate that search results accurately reflect the professor-course relationships established by administrators. Results must be complete and consistent with the database state.

Acceptance Criteria: Users can search by professor name and view all associated courses. Users can search by course name and view all assigned professors. The system displays clear, accurate, and complete results for each query. Search supports partial text matching and is case-insensitive. Friendly messages appear when no results are found. Search performance remains under two seconds for average query loads. The interface is responsive and adapts to desktop and mobile devices.

INVEST Characteristics: This story depends on US_CF_03, as accurate search requires established professor-course associations. Search filters, display format, and pagination are negotiable based on usability feedback. The story provides core functionality for navigating relationships, improving user experience significantly. Effort is estimable based on filter implementation, database queries, and frontend integration. The scope fits within a single sprint for a small development team. Testing validates that results accurately match linked data between professors and courses.

US_CourseFinder_05: Home Page and Access Flow for Administrator and User This story defines the main entry point of the system, providing role-appropriate navigation paths for administrators and general users.

User Need: Administrators want to access a dashboard summarizing key metrics with shortcuts to management actions. General users want immediate access to the search feature without unnecessary navigation.

Justification: Role-based entry points enable both user types to efficiently access the system according to their needs—administrators can manage data quickly, and users can immediately start searching.

Key Validations: The system must correctly identify the user’s role after login. Role-based routing must function reliably, directing administrators to the dashboard and general users to the search interface. The homepage must load within acceptable performance thresholds (under two seconds). Administrative elements must be completely hidden from general user views.

Acceptance Criteria: The system redirects users based on their role: administrators to the dashboard, general users to the search page. The homepage loads correctly and displays personalized content. The admin dashboard shows relevant metrics (total professors, total courses, recent updates) and quick management options. The general user view focuses exclusively on search functionality. The interface is responsive and loads within performance thresholds. Navigation between homepage and other sections works without errors.

INVEST Characteristics: This story depends on US_CF_04 and the authentication module, as it requires functional search capabilities and role identification. Homepage layout, design elements, and dashboard widgets are negotiable based on feedback and usability tests. The story provides personalized entry points, ensuring efficient access and navigation for both roles. The scope is measurable by estimating layout design, routing logic, and conditional rendering. Implementation fits within a single sprint. Testing verifies redirection logic, role-based rendering, and proper section access.

4.1.3 Story Management and Traceability

The user stories were managed in the Taiga project management platform, where they were assigned story points, prioritized in the product backlog, and allocated to specific


sprints during sprint planning sessions. Each story’s progress was tracked through states (New, In Progress, Done) and linked to specific tasks. Story points were estimated using planning poker sessions, with the team achieving consensus on complexity based on technical effort, uncertainty and integration requirements.

The dependency chain from administrative stories through search functionality to role-based navigation guided sprint planning, ensuring that foundational data management capabilities were delivered before user-facing features. This traceability ensures that every implemented feature can be validated against its original acceptance criteria and that test coverage aligns with user expectations. The acceptance criteria defined in each story directly influenced the test scenarios documented in Section 5 and the evaluation metrics presented in Section 6.

4.2 Web User Interface

This section presents the main artifacts developed during all the Workshops, which represent the foundation for the system implementation. Each artifact serves a specific purpose in documenting the design decisions and validating the feasibility of the proposed solution. The progression from conceptual models to concrete specifications ensures traceability from requirements to implementation.

Figure 1: Web User Interface (v1.0)



The image shows a web user interface for 'Academic Course Finder'. The main heading is 'Academic Course Finder' in a bold, dark blue font. Below this is a white, rounded rectangular box containing the 'Admin Login' section. The 'Admin Login' section has a title 'Admin Login' in bold dark blue. It includes two input fields: 'User' and 'Password', both with light blue borders and placeholder text. Below these fields is a dark blue 'Login' button with white text. Under the button is a link that says 'Recover your password' in a smaller, dark blue font. The entire form is centered on a light gray background. At the bottom of the page, there is a dark blue footer bar with the text '© 2025 Academic Course Finder. All rights reserved.' in white.

Academic Course Finder

Admin Login

User


Password

Login

[Recover your password](#)

© 2025 Academic Course Finder. All rights reserved.

Figure 2: Web User Interface (v1.0)



The image displays a web user interface for an "Academic Course Finder". The interface is contained within a light gray rectangular area. At the top of this area is a dark blue header bar with the text "Academic Course Finder" in white. Below the header, there is a search input field with the placeholder text "Search by Professor or course...". To the right of the input field is a dark blue button with the text "Search >". Below the search field is a large white rectangular box with the text "Search result" at the top left. At the bottom of the interface is a dark blue footer bar with the text "© 2025 Academic Course Finder. All rights reserved." in white.

Academic Course Finder

Search by Professor or course...

Search >

Search result

© 2025 Academic Course Finder. All rights reserved.

Figure 3: Web User Interface (v1.0)

Academic Course Finder

Search by Professor or course...

Search >

Professors teaching Calculus

Name: -----
Code: -----
Email: -----

Name: -----
Code: -----
Email: -----

Name: -----
Code: -----
Email: -----

© 2025 Academic Course Finder. All rights reserved.

Figure 4: Web User Interface (v1.0)



Figure 5: Web User Interface (v1.0)





Figure 6: Web User Interface (v1.0)



The user interface was designed following the principle "Simple. Clear. Consistent." The visual design prioritizes clarity, usability, and role-based access control to ensure that each stakeholder group can efficiently accomplish their tasks without unnecessary complexity. Six main views were designed and subsequently implemented in the production environment, representing the complete user journey through the system.

The interface accommodates three distinct user roles with different access levels and capabilities. Anonymous users can access the search functionality, allowing them to query courses by code, title, or professor details without authentication. The search results display comprehensive course information including schedules and associated instructors. Authenticated administrators access additional views for data management, including CRUD operations for courses and professors, as well as the assignment interface that links professors to specific courses and groups.

Design decisions emphasized consistency across views through standardized navigation patterns, uniform typography, and coherent color schemes that distinguish between public and administrative sections. The responsive layout adapts to different screen sizes, ensuring accessibility across desktop and mobile devices. Form validation provides immediate feedback to users, reducing errors during data entry. The role-based design ensures that administrative functions remain hidden from unauthorized users while maintaining an intuitive experience for all stakeholder groups.

The previous figures presents the final implemented views of the web application, demonstrating the translation from initial mockups to functional interfaces. These views validate the usability requirements established during the requirements analysis phase and confirm that the system meets the interaction needs identified in the user stories.

4.3 Architecture Overview

The system architecture implements a microservices approach with clear separation of concerns across three main components: frontend, an authentication service and a business logic service. This architectural decision was mandated by project requirements and offers several strategic advantages for development, deployment, and maintenance.

The separation of authentication into an independent service provides security isolation, ensuring that credential management and token generation occur in a dedicated, hardened environment. This service connects to a database optimized for transaction processing and user session management. The authentication service issues JWT tokens upon successful login, which are subsequently validated by the business logic service for

protected endpoints. This token-based approach enables stateless authentication, simplifying horizontal scaling and load balancing.

The business logic service handles all domain operations including course management, professor administration, search functionality, and relationship assignments. By isolating these operations in a separate service connected to a database different from the one used for authentication, the system achieves better performance through database specialization.

Component communication follows RESTful principles through well-defined API contracts. The frontend makes HTTPS requests to both backend services, with the authentication service providing tokens and the business service consuming them for authorization. This loose coupling allows independent deployment and version management of each service. Furthermore, the architecture supports parallel development, enabling team members to work simultaneously on different services without integration conflicts.




Figure 7: Architecture Diagram (v1.0).

Figure 7 illustrates the high-level architecture, showing the relationships between components, the data flow from user interaction through business logic to persistence layers, and the external interfaces exposed by each service. This diagram serves as the reference model for understanding system boundaries and component responsibilities throughout the development process.

4.4 Data Model and Database Schema

The data model defines four primary entities that capture the essential academic domain: Professor, Course, Admin, and ProfessorCourse. This structure was designed to support the core user requirements identified during the analysis phase, particularly the need to search courses by multiple criteria and maintain clear associations between professors and their teaching assignments.

4.4.1 Domain Entities and Relationships

The Professor entity contains identification attributes (id, name) and academic metadata (email, specialty) that categorizes the instructor's area of expertise. This information enables filtering and grouping operations in search queries. The email field is enforced as unique to prevent duplicate professor registrations and serves as a natural key for external integrations.

The Course entity includes structural attributes (id, code, name) along with scheduling information (schedule) that supports both course discovery and academic planning. The unique course code constraint ensures data integrity and prevents duplicate registrations. The schedule field stores time and day information in a flexible text format, allowing representation of various scheduling patterns without imposing rigid structural constraints.

The relationship between professors and courses implements a many-to-many association through the ProfessorCourse junction entity. This design decision reflects the reality of academic scheduling where one professor may teach multiple courses and one course may be taught by multiple professors across different groups or semesters. The ProfessorCourse entity includes additional context such as assigned_at (timestamp of assignment creation) and status (to manage active, inactive, pending, or completed assignments). This rich association model provides the flexibility required for complex academic scenarios while maintaining referential integrity through foreign key constraints with cascade delete behavior.

The Admin entity remains architecturally separate from the academic domain, containing only authentication-relevant attributes (id, username, password_hash, email). This separation ensures that administrative access control does not interfere with the core domain logic and allows for independent security policy management. The Admin entity resides in a separate database dedicated to the authentication service, while the academic entities (Professor, Course, ProfessorCourse) reside in a database managed by the business logic service.

Design decisions prioritized query performance for the primary use case: searching courses by professor or course attributes and retrieving all related information in minimal database operations. The normalized structure prevents data redundancy while the junction entity enables efficient joins. Repository interfaces abstract database operations, providing a clean separation between domain logic and persistence mechanisms that facilitates future database migration or optimization.

4.4.2 Business Logic Database

The business logic database implements the academic domain model. Table 2 summarizes the three main tables with their key attributes and constraints.

Table 2: Schema - Academic Domain Tables

Table	Key Attributes	Constraints
professors	id (SERIAL), name, email, specialty, created_at, updated_at	UNIQUE(email), Indexes on email, name, specialty
courses	id (SERIAL), name, code, schedule, created_at, updated_at	UNIQUE(code), Indexes on code, name, schedule
professor_courses	id (SERIAL), professor_id, course_id, assigned_at, status	FK to professors/courses, UNIQUE(professor_id, course_id), CHECK status IN ('active', 'inactive', 'pending', 'completed')

The **professors** table stores instructor information with automatic timestamp tracking. Indexes on email, name, and specialty optimize common query patterns including professor search by name and filtering by academic area. The email uniqueness constraint prevents duplicate registrations and supports using email as a natural identifier in external systems.

The **courses** table maintains course catalog information with indexed fields supporting fast lookup by course code or name. The schedule field stores time and location information in a flexible VARCHAR format, allowing various representations ("Mon/Wed 10:00-12:00", "Tue/Thu 14:00-15:30") without rigid parsing requirements. The unique code constraint ensures each course has a distinct identifier within the institution.

The `professor_courses` junction table implements the many-to-many relationship with additional assignment metadata. Foreign key constraints with CASCADE DELETE ensure referential integrity—deleting a professor or course automatically removes associated assignments, preventing orphaned records. The composite unique constraint on (`professor_id`, `course_id`) prevents duplicate assignments of the same professor to the same course. The status field uses a CHECK constraint limiting values to predefined states, supporting assignment lifecycle management. Multiple indexes optimize queries retrieving courses by professor, professors by course, and filtering by assignment status. A composite index on (`professor_id`, `course_id`, `status`) accelerates the most common query pattern: finding active courses for a specific professor.

Database triggers automatically update the `updated_at` timestamp on all tables whenever records are modified, providing audit trail capabilities without application-level logic. This pattern ensures consistent timestamp management across all domain entities.

4.4.3 Authentication Database

The authentication service uses a separate database (`auth_db_course`) to isolate security-critical data from business logic. Table 3 describes the authentication schema structure.

Table 3: Schema - Authentication Table

Table	Key Attributes	Constraints
<code>admins</code>	<code>id</code> (INT AUTO_INCREMENT), <code>username</code> , <code>password_hash</code> , <code>email</code> , <code>created_at</code> , <code>updated_at</code>	UNIQUE(<code>username</code>), UNIQUE(<code>email</code>), Indexes on <code>username</code> and <code>email</code>

The `admins` table stores administrative user credentials with bcrypt-hashed passwords in the `password_hash` field. Both `username` and `email` are enforced as unique, providing multiple authentication options. Indexes on both fields optimize login queries, which typically search by `username`.

The deliberate separation of authentication and business databases supports the microservices architecture by allowing independent scaling, backup strategies, and security policies. The authentication database requires higher security measures while the business database optimizes for read-heavy query patterns typical of search operations.

4.4.4 Schema Design Rationale

Several design decisions warrant explicit justification. First, the choice of different database systems (PostgreSQL for business logic, MySQL for authentication) was mandated by project requirements but offers strategic advantages. PostgreSQL’s advanced JSON support, full-text search capabilities, and complex query optimization make it well-suited for the search-intensive academic domain. MySQL’s simpler deployment and mature authentication integrations align with the lightweight requirements of the authentication service.

Second, the denormalized storage of schedule information as text rather than structured time slots prioritizes implementation simplicity and flexibility over query complexity. While this approach limits schedule-based filtering capabilities, it satisfies current requirements and can be refactored if advanced schedule analysis becomes necessary.

Third, the status field in `professor_courses` supports future features like temporary assignments, historical tracking, and workflow management without schema changes. The CHECK constraint ensures data quality while allowing extension through configuration rather than migration.

Finally, comprehensive indexing strategies reflect analysis of expected query patterns: searches by professor name, searches by course code/name, and retrieval of professor-course relationships. Composite indexes optimize multi-condition queries while single-column indexes support individual lookups. This indexing approach balances query performance against write overhead and storage costs.




Figure 8: Class diagram showing domain entities, repositories, services, and controllers across all architectural layers.

Figure 8 presents the complete class diagram including entities, repositories, services, and controllers across all architectural layers. This diagram demonstrates how the data model integrates with business logic components and exposes functionality through controller endpoints. The layered structure visible in the diagram reinforces the architectural principles of separation of concerns and dependency management established in the sys-

tem design. The repository interfaces shown in the diagram directly correspond to the database tables documented in this section, providing the abstraction layer that enables database-agnostic business logic implementation.

4.5 API Endpoints and Contracts

The system exposes its functionality through two RESTful API services that communicate using JSON payloads over HTTP/HTTPS protocols. Each service provides a well-defined set of endpoints documented through OpenAPI specifications and accessible via Swagger UI for interactive exploration and testing. This section describes the API contracts, request/response formats, and authentication mechanisms that enable communication between the frontend and backend services.

4.5.1 Authentication Service API

The authentication service, built with Java Spring Boot and running on port 8080, provides JWT-based authentication for administrative users. The API documentation is accessible through Swagger UI at <http://localhost:8080/swagger-ui/index.html> during development, offering an interactive interface to explore endpoints, test requests, and view response formats.

Login Endpoint The primary authentication endpoint validates administrator credentials and issues JWT tokens for subsequent authorized requests.

Endpoint: POST /api/auth/login

Description: Authenticates an administrator by validating username and password against stored credentials in the MySQL database. Upon successful authentication, the service generates a JWT token containing the admin's identity and role claims.

Figure 9 shows the Swagger UI documentation for the authentication service, illustrating the login endpoint specification with its request schema and expected responses. The interactive interface allows developers to test authentication directly from the browser by providing sample credentials and observing the token generation process.

Figure 10 demonstrates the request format for the login endpoint. The request body requires two fields: username and password, both as string values. The Swagger interface validates the request schema before submission, ensuring that required fields are present and properly formatted.




Figure 9: Swagger UI documentation for the authentication service showing the login endpoint specification with request/response schemas.




Figure 10: Login endpoint request example showing the JSON payload structure with username and password fields.

Figure 11 shows the successful response format. Upon valid credentials, the service returns a JSON object containing a JWT token string. This token must be included in the Authorization header of subsequent requests to protected endpoints in both services. The token format follows the Bearer authentication scheme: **Authorization: Bearer <token>**. Tokens expire after a configured period (typically 24 hours) and must be refreshed through re-authentication.




Figure 11: Successful login response showing the generated JWT token that must be used for authenticated requests.

Error responses return appropriate HTTP status codes with descriptive messages. A 401 Unauthorized status indicates invalid credentials, while 400 Bad Request indicates malformed request payloads missing required fields.

4.5.2 Business Logic Service API

The business logic service, built with Python FastAPI and running on port 8000, provides comprehensive CRUD operations for professors, courses, and their assignments. All administrative endpoints require valid JWT tokens obtained from the authentication service. The API documentation is accessible at <http://localhost:8000/docs> during development.

The API follows RESTful conventions with consistent HTTP method semantics: POST for creation, GET for retrieval, PUT for updates, and DELETE for removal. All endpoints return JSON responses with appropriate HTTP status codes (200 for success, 201 for creation, 400 for validation errors, 401 for authentication failures, 404 for not found, 500 for server errors).

Figure 12 presents the comprehensive Swagger UI documentation for the business logic

service, showing all available endpoints organized by resource type (professors, courses, assignments). The interface provides expandable sections for each endpoint with detailed parameter descriptions, request schemas, and response models.




Figure 12: Swagger UI overview of the business logic service showing all CRUD endpoints organized by resource (professors, courses, assignments).

Professor Endpoints The professor endpoints manage instructor data and support the administrative user stories related to professor registration and management.

- **POST /api/v1/professors** — Create a new professor. Requires authentication. Validates email uniqueness and required fields (name, email, specialty). Returns 201 Created with the created professor object including the assigned ID.
- **GET /api/v1/professors** — List all professors with pagination support. Accepts query parameters: **page** (default: 1), **limit** (default: 10), **search** (optional text filter). Returns paginated results with total count and page metadata.
- **GET /api/v1/professors/{id}** — Retrieve a specific professor by ID. Returns 404 if professor not found. Response includes all professor attributes and a list of assigned courses.
- **PUT /api/v1/professors/{id}** — Update professor information. Requires authentication. Validates email uniqueness excluding the current professor. Returns updated professor object.
- **DELETE /api/v1/professors/{id}** — Delete professor. Requires authentication. Cascade deletes all professor-course assignments due to foreign key constraints. Returns 204 No Content on success.

Figure 13 shows the Swagger documentation for the professor creation endpoint. The interface displays the required request body schema with field descriptions, data types, and validation constraints. The example section provides a sample payload that developers can modify and submit directly through the Swagger UI for testing purposes.




Figure 13: Professor creation endpoint documentation showing request schema with required fields (name, email, specialty) and expected response format.

Course Endpoints The course endpoints manage academic course catalog data, supporting course registration and discovery functionality.

- **POST /api/v1/courses** — Create a new course. Requires authentication. Validates code uniqueness and required fields (name, code). Schedule is optional. Returns 201 Created with the created course object.
- **GET /api/v1/courses** — List all courses with pagination and search capabilities. Accepts query parameters: `page`, `limit`, `search` (filters by course name or code). Returns paginated results.
- **GET /api/v1/courses/{id}** — Retrieve specific course by ID. Returns course details including all assigned professors with their assignment status.

- **PUT** `/api/v1/courses/{id}` — Update course information. Requires authentication. Validates code uniqueness excluding the current course. Returns updated course object.
- **DELETE** `/api/v1/courses/{id}` — Delete course. Requires authentication. Cascade deletes all professor-course assignments. Returns 204 No Content on success.

Figure 14 illustrates the complete set of course endpoints available in the Swagger documentation. Each endpoint section can be expanded to reveal detailed parameter information, authentication requirements, and example request/response payloads.

Courses		^
GET	<code>/api/v1/courses</code> Get all courses	⌵
POST	<code>/api/v1/courses</code> Create a new course	🔒 ⌵
GET	<code>/api/v1/courses/paginated</code> Get courses paginated	⌵
GET	<code>/api/v1/courses/{course_id}</code> Get course by ID	⌵
PUT	<code>/api/v1/courses/{course_id}</code> Update course	🔒 ⌵
DELETE	<code>/api/v1/courses/{course_id}</code> Delete course	🔒 ⌵

Figure 14: Course CRUD endpoints documentation showing all available operations (POST, GET, PUT, DELETE) with their respective paths and authentication requirements.

Assignment Endpoints The assignment endpoints manage the many-to-many relationships between professors and courses, implementing the core linking functionality described in user story US_CF_03.

- **POST** `/api/v1/assignments` — Create a professor-course assignment. Requires authentication. Validates that both professor and course exist. Enforces unique constraint preventing duplicate assignments. Returns 201 Created with assignment details.
- **GET** `/api/v1/assignments` — List all assignments with optional filtering by professor ID or course ID. Supports pagination. Returns assignments with embedded professor and course information.
- **GET** `/api/v1/assignments/{id}` — Retrieve specific assignment by ID. Returns assignment details including full professor and course objects.
- **PUT** `/api/v1/assignments/{id}` — Update assignment status. Requires authentication. Allows changing status between active, inactive, pending, and completed. Returns updated assignment.
- **DELETE** `/api/v1/assignments/{id}` — Remove professor-course assignment. Requires authentication. Returns 204 No Content on success.

Figure 15 displays the assignment creation endpoint documentation. The request schema shows the required fields (`professor_id`, `course_id`, `status`) and their data types. The response schema includes the complete assignment object with embedded professor and course details, demonstrating the API's approach to reducing client-side requests by returning related data in a single response.

POST /api/v1/assignments Create assignment

Parameters
No parameters

Request body required
application/json

Example Value **Schema**

```
{
  "professor": {
    "name": "string",
    "email": "user@example.com",
    "specialty": "string",
    "id": 0
  },
  "course": {
    "name": "string",
    "code": "string",
    "schedule": "string",
    "id": 0
  },
  "status": "active"
}
```

Responses

Code	Description	Links
201	Successful Response	No links
422	Validation Error	No links

201 Successful Response
Media type: application/json
Controls Accept header

Example Value **Schema**

```
{
  "professor_id": 0,
  "course_id": 0,
  "status": "active",
  "id": 0,
  "assigned_at": "2025-12-11T21:28:35.467Z",
  "professor": {
    "name": "string",
    "email": "user@example.com",
    "specialty": "string",
    "id": 0
  },
  "course": {
    "name": "string",
    "code": "string",
    "schedule": "string",
    "id": 0
  }
}
```

422 Validation Error
Media type: application/json

Example Value **Schema**

```
{
  "detail": [
    {
      "loc": [
        "string",
        0
      ],
      "msg": "string",
      "type": "string"
    }
  ]
}
```

Figure 15: Assignment creation endpoint showing request schema with `professor_id`, `course_id`, and `status` fields, along with the nested response structure including professor and course details.

Search Endpoints While not explicitly listed in the core CRUD operations, the business logic service implements specialized search endpoints that support user story US_CF_04. These endpoints optimize query performance through database indexes and implement the bidirectional search functionality.

- **GET /api/v1/search/courses-by-professor?name=<professor_name>** — Search courses taught by a specific professor. Supports partial name matching and case-insensitive search. Returns list of courses with professor details.

- GET `/api/v1/search/professors-by-course?name=<course_name>` — Search professors teaching a specific course. Supports partial course name or code matching. Returns list of professors assigned to matching courses.

4.5.3 API Design Principles and Error Handling

Both services follow consistent design principles that enhance usability and maintainability. All endpoints return structured JSON responses with predictable formats. Error responses include human-readable messages, error codes, and timestamps to facilitate debugging. Validation errors return 400 Bad Request with detailed field-level error descriptions specifying which fields failed validation and why.

Authentication failures return 401 Unauthorized with clear indication that credentials are invalid or tokens have expired. Authorization failures (attempting to access admin endpoints without proper tokens) return 403 Forbidden. Resource not found scenarios return 404 Not Found with the resource type and ID that could not be located.

Pagination follows cursor-based patterns with consistent query parameters across all list endpoints. Responses include metadata about total items, current page, and page size, enabling frontend components to implement efficient pagination controls. Search endpoints support partial matching and case-insensitive queries to improve user experience and reduce friction during course discovery.

The API versioning strategy uses URL path prefixes (`/api/v1/`) to enable future backward-compatible changes without breaking existing integrations. This approach allows the introduction of new API versions while maintaining support for legacy clients during transition periods.

CORS (Cross-Origin Resource Sharing) configuration allows the React frontend to communicate with both backend services during development and production. The authentication service validates JWT tokens issued by itself, while the business logic service validates tokens by verifying signatures using the shared secret key configured during deployment.

Table 4 summarizes the complete API surface across both services, providing a reference for the total number of endpoints and their authentication requirements.

The API design supports the implementation of all user stories documented in Section 4 while providing the flexibility needed for future feature additions. The consistent patterns across endpoints reduce frontend development complexity and enable code reuse

Table 4: API Endpoints Summary

Service	Endpoints	Auth Required	Primary Purpose
Authentication	1	No	JWT token generation
Professors	5	Yes (except GET list/single)	Professor CRUD
Courses	5	Yes (except GET list/single)	Course CRUD
Assignments	5	Yes (except GET list)	Assignment management
Search	2	No	Bidirectional search
Total	18		

through shared API client libraries.

4.6 Deployment View




Figure 16: Deployment diagram (v1.0).

5 Implementation and Quality Assurance

5.1 Testing Strategy

The testing strategy for the Course Finder platform follows a multi-layered approach designed to validate functionality at different levels of the system architecture. Tests were implemented after completing the initial development phase, focusing on critical business logic components, API endpoints, user story acceptance criteria, and system performance under load. The overall philosophy prioritizes thorough coverage of core functionality and user-facing features while ensuring that the system meets the quality expectations defined in the acceptance criteria of each user story.

The testing approach focused on meaningful validation of critical paths: unit tests for business logic components related to API endpoints, acceptance tests mapping directly to user stories, and stress tests evaluating system performance and scalability. This pragmatic strategy ensures that tests provide real value in detecting defects and validating requirements.

5.1.1 Unit Testing

Unit tests validate individual components in isolation, ensuring that business logic functions correctly independent of external dependencies such as databases or network services. Mock objects simulate repository and service interactions, allowing precise control over test scenarios and enabling rapid test execution without infrastructure dependencies.

Python Backend - Business Logic Service The business logic service implemented includes comprehensive unit tests using pytest. Tests cover the core service layer components responsible for CRUD operations on professors, courses, and assignments. The test suite uses fixtures to provide consistent test data and mock repository objects that simulate database operations without requiring actual PostgreSQL connections.

Table 5 summarizes the unit test coverage for the Python backend services. Each service class includes dedicated test classes organized by operation type (Create, Read, Update, Delete), following the Arrange-Act-Assert pattern for clarity and maintainability.

The ProfessorService tests validate critical scenarios including successful professor creation with unique email addresses, prevention of duplicate email registration through HTTPException with 409 Conflict status, retrieval of individual professors and paginated lists, handling of non-existent professor queries with 404 Not Found responses, successful updates with field modifications, and cascade deletion behavior. Each test uses mock

Table 5: Python Backend Unit Tests Summary

Service	Tests	Scenarios Covered
ProfessorService	13	Create (success, duplicate email), Read (by ID, all, paginated, not found), Update (success, not found, duplicate email), Delete (success, not found)
CourseService	14	Create (success, duplicate code), Read (by ID, all, paginated, empty list, not found), Update (success, not found, duplicate code), Delete (success, not found)
Total	27	

repository objects that return predefined data, allowing precise verification of service layer logic without database dependencies.

Similarly, the CourseService tests cover analogous scenarios for course entities, with particular emphasis on course code uniqueness validation and schedule information handling. Tests verify that the service correctly enforces business rules such as preventing duplicate course codes, returning appropriate error responses for invalid operations, and properly transforming domain models into response schemas.

Mock fixtures defined in `conftest.py` provide reusable test data including sample professor and course objects with realistic attributes, schema instances for create and update operations, and pre-configured mock repository objects with all necessary methods. This fixture-based approach reduces test code duplication and ensures consistency across test cases.

Java Backend - Authentication Service The authentication service implemented includes unit tests using JUnit 5 to validate authentication logic and JWT token generation. Tests focus on two critical components: the authentication service responsible for credential validation and the JWT utility class responsible for token generation and validation.

The AuthService tests cover three primary scenarios using mocked dependencies. First, successful login tests verify that valid username and password combinations result in JWT token generation, with the AdminRepository mock returning a valid admin

entity and the PasswordEncoder mock confirming password match. Second, failed login due to incorrect password tests verify that the service throws appropriate authentication exceptions when the password encoder indicates a mismatch, even when the user exists in the repository. Third, failed login due to non-existent user tests validate that queries for unknown usernames result in authentication failures with appropriate error messages.

The JwtUtil tests verify token generation correctness by creating tokens with test data and validating the resulting JWT structure. Tests confirm that generated tokens are non-null, contain the expected username as the subject claim, and include valid expiration timestamps set according to the configured token lifetime. These tests use temporary RSA key pairs instead of production keys, ensuring security while maintaining test validity.

Both test suites use extensive mocking to isolate the components under test from external dependencies. The AdminRepository is mocked to return predefined admin entities without database queries. The PasswordEncoder is mocked to control password matching behavior deterministically. The JwtUtil mock (when testing AuthService) provides predictable token values. This isolation ensures that tests execute quickly and reliably regardless of external system state.

5.1.2 Acceptance Testing with Cucumber

Acceptance tests validate that the implemented system satisfies the user stories and acceptance criteria defined during requirements analysis. These tests use Cucumber to express scenarios in Gherkin syntax, providing readable specifications that map directly to user story acceptance criteria and can be executed as automated tests.

Four feature files cover the core user stories documented in Section 4. Each feature includes multiple scenarios that validate both successful operations and error handling for edge cases.

Feature: Registration of New Teachers (US_CF_01) This feature validates professor registration functionality through four scenarios. The first scenario verifies that administrators can access the registration form from the main panel. The second scenario validates successful professor creation when all required fields (name, email, specialty) are provided with valid data, confirming that the system displays success messages and updates the professor list without page reload. The third scenario tests validation behavior when required fields are left empty, ensuring appropriate error messages guide administrators to correct the input. The fourth scenario verifies duplicate prevention by attempting to register a professor with an email or ID that already exists, confirming

that the system prevents creation and displays clear error messages.

Feature: Registration of New Subjects (US_CF_02) This feature validates course registration through five scenarios covering similar patterns to professor registration. Scenarios test form access, successful registration with valid data, required field validation, duplicate course prevention based on course code, and optional teacher assignment during course creation. The teacher assignment scenario specifically validates that when an administrator assigns an existing teacher during course creation, the relationship is correctly stored in the database and visible in both professor and course views.

Feature: Assignment of Professors to Courses (US_CF_03) This feature addresses the core relationship management functionality through seven scenarios. Initial scenarios validate interface access and successful assignment creation with valid professor and course selections. Validation scenarios ensure that missing required fields prevent assignment creation with appropriate error messages. Two scenarios specifically test the many-to-many relationship behavior: one verifying that a single professor can be assigned to multiple courses, and another confirming that a single course can have multiple professors assigned. A duplicate prevention scenario validates that the system enforces the unique constraint on professor-course pairs, preventing redundant assignments. Finally, a deletion scenario confirms that administrators can remove assignments and that the system correctly updates the interface.

Feature: Search Courses by Professor and Professors by Course (US_CF_04) This feature validates the primary end-user search functionality through four scenarios. The first scenario tests searching courses by professor name, verifying that all courses taught by the specified professor appear in results. The second scenario validates the inverse operation: searching professors by course name. The third scenario specifically tests partial matching and case-insensitive search behavior, ensuring that users can find results even with incomplete or lowercase search terms. The fourth scenario validates graceful handling of searches with no matches, confirming that the system displays friendly "No results found" messages rather than errors or empty pages.

Table 6 summarizes the acceptance test coverage across all features, showing the direct mapping from user stories to automated test scenarios.

All 20 acceptance test scenarios pass successfully, confirming that the implemented

Table 6: Cucumber Acceptance Test Scenarios Summary

Feature / User Story	Scenarios	Status
US_CF_01: Registration of Teachers	4	Passing
US_CF_02: Registration of Subjects	5	Passing
US_CF_03: Assignment Management	7	Passing
US_CF_04: Bidirectional Search	4	Passing
Total	20	All Passing

system satisfies the requirements documented in the user stories. The Gherkin scenarios serve dual purposes: they provide executable specifications that validate system behavior automatically, and they function as living documentation that stakeholders can read to understand system capabilities without technical expertise.

The acceptance tests execute against the full application stack including frontend, both backend services, and test databases populated with seed data. This end-to-end validation ensures that component integration functions correctly and that the user experience matches specifications. Test execution occurs both during local development and as part of the CI/CD pipeline, providing continuous validation that new changes do not break existing functionality.

5.1.3 Stress Testing with JMeter

Stress testing evaluates system performance and scalability under realistic and peak load conditions. These tests identify performance bottlenecks, validate response time requirements specified in user stories, and ensure that the system remains stable and responsive when serving multiple concurrent users. Stress tests using Apache JMeter are currently under active development and will be completed in the next project phase.

The planned stress testing strategy will focus on critical endpoints identified through user story analysis. Primary targets include the authentication endpoint (login), search endpoints (courses by professor, professors by course), and administrative CRUD operations (professor creation, course creation, assignment management). Test scenarios will simulate realistic usage patterns with varying levels of concurrency to identify system behavior under normal load, peak load, and stress conditions.

Initial test plans specify scenarios with 50, 100, and 500 concurrent users executing typical workflows over sustained periods (5-10 minutes). Metrics to be collected include average response time, 95th percentile response time, throughput (requests per second), error rate, and resource utilization (CPU, memory, database connections). These met-

rics will validate whether the system meets the performance requirement specified in US_CF_04 and US_CF_05 (response times under two seconds for average query loads).

Results from stress testing will inform infrastructure sizing decisions for production deployment, identify optimization opportunities in database queries or API implementations, and validate that the system architecture can scale to support the anticipated user base. Performance baselines established through these tests will guide future capacity planning and serve as regression benchmarks for subsequent releases.

5.1.4 Testing Infrastructure and Automation

The testing infrastructure integrates with the CI/CD pipeline implemented through GitHub Actions. Unit tests execute automatically on every commit and pull request, providing immediate feedback to developers about test failures or regressions.

Acceptance tests execute as part of the integration testing phase after successful unit test completion. These tests require more setup time as they spin up Docker containers for both backend services and test databases, populate seed data, and initialize the frontend application. The complete acceptance test suite executes in approximately 5-10 minutes, providing thorough validation before deployment to staging environments.

Test results are reported through the GitHub Actions interface, with detailed logs available for investigation when failures occur. Failed tests block pull request merges, ensuring that defects are identified and resolved before integration into the main codebase. This continuous validation approach maintains code quality and prevents regression of previously working functionality.

Local development environments support test execution through simple commands (`pytest` for Python unit tests, `mvn test` for Java unit tests, `npm run test:cucumber` for acceptance tests), enabling developers to validate changes before committing. Mock-based unit tests execute extremely quickly in local environments, supporting test-driven development workflows even for components with complex dependencies.

5.2 CI/CD Pipeline

5.3 Code Quality and Standards

6 Results and Discussion

The outcomes of Workshops 1 and 2 have established a solid analytical foundation for the project. The main deliverables include the problem definition, user stories with acceptance criteria, CRC cards, architectural and deployment diagrams, and early Web UI mockups. These artifacts provide a coherent view of the system's objectives and structure, enabling the team to validate the feasibility of the proposed solution before implementation.

6.1 Testing Results

6.2 System Performance Evaluation

6.3 Achievement of Objectives

6.4 Limitations and Lessons Learned

7 Conclusions

This first report consolidates the analytical and design foundations of the project, demonstrating coherent progress from requirement elicitation to system architecture definition. The established artifacts and planning ensure a clear path toward implementation, testing, and integration in the following workshop stages.

Glossary

API Application Programming Interface.

CI/CD Continuous Integration / Continuous Deployment.

JWT JSON Web Token.

CRC Class - Responsibility - Collaboration Cards